


Introduction to Natural Language processing

Natural language processing (NLP): Subfield of linguistics, computer science, and information engineering concerned with the interactions between computers and human (natural) languages, in particular how to program computers to process and analyze large amounts of natural language data.



What is it all about?

What's that all about

Let's start with the basics:



we speak



we read



we write

all that using language

But that's not all...



- we think of the world around us
- we dream
- we make decisions and plans
- all in natural language, i.e. in words

What are the raw materials for NLP revolution?

1. Data: There is a huge amount of human knowledge that has been digitized.

2. Compute: Our personal computers are very powerful, and it has never been easier to rent more compute power.


3. Algorithms: We use a combination of Computer Science, Machine Learning, Deep Learning, and linguist techniques.

Why NLP is hard?

It's the nature of the human language that makes NLP difficult.

Some examples:

1. Ambiguity: One word/sentence can have different meanings.



2. Spelling Errors: There are some words that are usually misspelled by people.



Top 10 Misspelled Words

separate ✓

seperate ✗

calendar ✓

calender ✗

unnecessary ✓

unecessary ✗

questionnaire ✓

questionaire ✗

accidentally ✓

accidentaly ✗

necessary ✓

neccessary ✗

surprise ✓

suprise ✗

embarrass ✓

embarass ✗

accommodation ✓

acomodation ✗

receive ✓

recieve ✗

3. Phonetics and Phonology: Different words that have the same pronunciation.

- there / their
- plane / plain
- ice cream / I scream

4. Order of Context: "Dog bites man" and "Man bites dog" have vastly different meanings

In ML terms, why is NLP hard?

1. Input: Large, discrete, open-ended input state spaces

2. Output: Variable length output spaces

3. Algorithms: Recycled from other fields

Top Used Text File Types

- CSV
- Text Files
- JSON
- SQL
- HTML
- PDF
- DOCX

NLP Libraries

There are many popular NLP libraries that we will cover in the next labs. In this lab, we will get introduced to NLTK.

NLTK

Natural Language Toolkit (nltk) is a Python package for NLP <https://www.nltk.org/>

What are Corpus / Corpora?

a collection of written texts, especially the entire works of a particular author or a body of writing on a particular subject.

```
In [1]: !pip install nltk
```

Collecting nltk

Downloading nltk-3.10.3-py3-none-any.whl.metadata (3.2 kB)

Requirement already satisfied: defusedxml in c:\users\huawei\miniconda3\lib\site-packages (from nltk) (0.7.1)

Requirement already satisfied: click in c:\users\huawei\miniconda3\lib\site-packages (from nltk) (8.2.1)

Requirement already satisfied: joblib in c:\users\huawei\miniconda3\lib\site-packages (from nltk) (1.5.3)

Collecting regex<=2021.8.3 (from nltk)

Downloading regex-2026.9.10-cp313-cp313-win_amd64.whl.metadata (41 kB)

Requirement already satisfied: tqdm in c:\users\huawei\miniconda3\lib\site-packages (from nltk) (4.67.1)

Requirement already satisfied: colorama in c:\users\huawei\miniconda3\lib\site-packages (from click->nltk) (0.4.6)

Downloading nltk-3.10.3-py3-none-any.whl (1.8 MB)

```
----- 0.0/1.8 MB ? eta -:--:--
----- 0.0/1.8 MB ? eta -:--:--
----- 0.5/1.8 MB 2.7 MB/s eta 0:00:01
----- 1.3/1.8 MB 3.2 MB/s eta 0:00:01
----- 1.8/1.8 MB 3.2 MB/s eta 0:00:00
```

Downloading regex-2026.9.10-cp313-cp313-win_amd64.whl (277 kB)

Installing collected packages: regex, nltk

```
----- 0/2 [regex]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 1/2 [nltk]
----- 2/2 [nltk]
```

Successfully installed nltk-3.10.3 regex-2026.9.10

In [2]: `import nltk`

Arabic Libraries (Extra)

Camel-Tools: <https://camel-tools.readthedocs.io/en/latest/>

PyArabic: <https://pyarabic.readthedocs.io/ar/latest/>

Farasa: <https://farasa.qcri.org/>

NLP Applications

1. Information Extraction

The automated retrieval of specific information related to a selected topic from a body or bodies of text.


2. Machine Translation

The process of using artificial intelligence to automatically translate text from one language to another without human involvement.

3. Sentiment Analysis

Identifying sentiments and opinions stated in a text.

Example of an NLP Pipeline



1. Text Acquisition:

Obtain the text data from a source, such as a file, a web page, or a database.

2. Text Preprocessing:

- **Tokenization:** Split the text into individual words or tokens.
- **Lowercasing:** Convert all tokens to lowercase to ensure uniformity.
- **Removing Punctuation:** Strip away punctuation marks from tokens.

- **Stopword Removal:** Eliminate common words (e.g., "and," "the," "is") that may not carry significant meaning.
- **Stemming:** Removes prefixes/suffixes to get a root form. (Running -> run)
- **Lemmatization:** Converts word to base/dictionary form (lemma). (better -> good).

3. Feature Extraction:

- **Bag-of-Words (BoW):** Create a vector representation of the text using the frequency of each token in the document.
- **TF-IDF (Term Frequency-Inverse Document Frequency):** Assign weights to tokens based on their importance in the document relative to the entire corpus.

4. Model Building:

Choose a machine learning model for sentiment analysis, such as a Naive Bayes classifier, Support Vector Machine (SVM), or a neural network.

Task#1:

Students should work in groups and search online for one of the different types of datasets. • CSV • Text Files • JSON • SQL • HTML • PDF • DOCX

Task#2:

Download and open the dataset.

```
In [3]: conda install kagglehub
```

3 channel Terms of Service accepted
 Retrieving notices: done
 Channels:
 - defaults
 Platform: win-64
 Collecting package metadata (repodata.json): done
 Solving environment: done

Package Plan

environment location: c:\Users\huawei\miniconda3

added / updated specs:
 - kagglehub

The following packages will be downloaded:

package	build	
-----	-----	
ca-certificates-2026.8.13	haa95532_0	107 KB
certifi-2026.7.22	py313haa95532_0	135 KB
conda-26.1.1	py313haa95532_0	1.2 MB
kagglehub-1.0.2	py313haa95532_0	163 KB
kagglesdk-0.1.35	py313haa95532_0	438 KB
openssl-3.5.8	hbb43b14_0	8.9 MB
protobuf-7.35.1	py313h854c0a0_0	481 KB
-----	-----	
Total:		11.4 MB

The following NEW packages will be INSTALLED:

kagglehub	pkgs/main/win-64::kagglehub-1.0.2-py313haa95532_0
kagglesdk	pkgs/main/win-64::kagglesdk-0.1.35-py313haa95532_0
protobuf	pkgs/main/win-64::protobuf-7.35.1-py313h854c0a0_0

The following packages will be UPDATED:

ca-certificates	2025.7.15-haa95532_0 --> 2026.8.13-haa95532_0
certifi	2025.8.3-py313haa95532_0 --> 2026.7.22-py313haa95532_0
conda	25.7.0-py313haa95532_0 --> 26.1.1-py313haa95532_0
openssl	3.0.17-h35632f6_0 --> 3.5.8-hbb43b14_0

Downloading and Extracting Packages: ...working...

openssl-3.5.8	8.9 MB		0%
conda-26.1.1	1.2 MB		0%
protobuf-7.35.1	481 KB		0%
kagglesdk-0.1.35	438 KB		0%

kagglehub-1.0.2	163 KB		0%
certifi-2026.7.22	135 KB		0%
ca-certificates-2026	107 KB		0%
protobuf-7.35.1	481 KB	3	3%
kagglehub-1.0.2	163 KB	9	10%
protobuf-7.35.1	481 KB	#3	13%
kagglesdk-0.1.35	438 KB	3	4%
kagglehub-1.0.2	163 KB	###9	39%
openssl-3.5.8	8.9 MB		0%
protobuf-7.35.1	481 KB	##6	27%
kagglehub-1.0.2	163 KB	#####8	79%
kagglesdk-0.1.35	438 KB	#8	18%
kagglehub-1.0.2	163 KB	#####	100%
openssl-3.5.8	8.9 MB		1%
protobuf-7.35.1	481 KB	####3	43%

kagglesdk-0.1.35	438 KB	###2	33%
openssl-3.5.8	8.9 MB	1	1%
conda-26.1.1	1.2 MB	1	1%

kagglehub-1.0.2	163 KB	#####	100%
-----------------	--------	-------	------

certifi-2026.7.22	135 KB	#1	12%
-------------------	--------	----	-----

protobuf-7.35.1	481 KB	#####6	67%
-----------------	--------	--------	-----

kagglesdk-0.1.35	438 KB	#####2	62%
openssl-3.5.8	8.9 MB	1	2%
conda-26.1.1	1.2 MB	5	5%

certifi-2026.7.22	135 KB	#####	100%
-------------------	--------	-------	------

protobuf-7.35.1	481 KB	#####	100%
-----------------	--------	-------	------

kagglesdk-0.1.35	438 KB	#####	100%
openssl-3.5.8	8.9 MB	2	3%
conda-26.1.1	1.2 MB	#	10%

certifi-2026.7.22	135 KB	#####	100%
-------------------	--------	-------	------

certifi-2026.7.22	135 KB	#####	100%
openssl-3.5.8	8.9 MB	3	4%
conda-26.1.1	1.2 MB	#9	19%

protobuf-7.35.1	481 KB	#####	100%
protobuf-7.35.1	481 KB	#####	100%
openssl-3.5.8	8.9 MB	5	5%
conda-26.1.1	1.2 MB	##4	24%
ca-certificates-2026	107 KB	#4	15%
ca-certificates-2026	107 KB	#####	100%
ca-certificates-2026	107 KB	#####	100%
openssl-3.5.8	8.9 MB	8	9%
conda-26.1.1	1.2 MB	###7	37%
kagglesdk-0.1.35	438 KB	#####	100%
kagglesdk-0.1.35	438 KB	#####	100%
conda-26.1.1	1.2 MB	#####1	51%
openssl-3.5.8	8.9 MB	#3	13%
openssl-3.5.8	8.9 MB	#8	19%
conda-26.1.1	1.2 MB	#####5	76%
conda-26.1.1	1.2 MB	#####9	90%
conda-26.1.1	1.2 MB	#####	100%
openssl-3.5.8	8.9 MB	##1	21%
openssl-3.5.8	8.9 MB	##3	23%
openssl-3.5.8	8.9 MB	##5	25%
openssl-3.5.8	8.9 MB	##7	28%
openssl-3.5.8	8.9 MB	##9	30%
openssl-3.5.8	8.9 MB	###1	32%
openssl-3.5.8	8.9 MB	###4	34%
openssl-3.5.8	8.9 MB	###6	36%
openssl-3.5.8	8.9 MB	###9	39%
conda-26.1.1	1.2 MB	#####	100%

openssl-3.5.8	8.9 MB	#####1	42%
openssl-3.5.8	8.9 MB	#####4	44%
openssl-3.5.8	8.9 MB	#####6	47%
openssl-3.5.8	8.9 MB	#####9	49%
openssl-3.5.8	8.9 MB	#####1	52%
openssl-3.5.8	8.9 MB	#####4	55%
openssl-3.5.8	8.9 MB	#####7	57%
openssl-3.5.8	8.9 MB	#####9	60%
openssl-3.5.8	8.9 MB	#####2	63%
openssl-3.5.8	8.9 MB	#####5	66%
openssl-3.5.8	8.9 MB	#####7	68%
openssl-3.5.8	8.9 MB	#####1	71%
openssl-3.5.8	8.9 MB	#####4	74%
openssl-3.5.8	8.9 MB	#####7	77%
openssl-3.5.8	8.9 MB	#####9	80%
openssl-3.5.8	8.9 MB	#####3	84%
openssl-3.5.8	8.9 MB	#####6	86%
openssl-3.5.8	8.9 MB	#####	91%
openssl-3.5.8	8.9 MB	#####3	94%
openssl-3.5.8	8.9 MB	#####7	97%
openssl-3.5.8	8.9 MB	#####	100%
openssl-3.5.8	8.9 MB	#####	100%

done

Preparing transaction: done
 Verifying transaction: done
 Executing transaction: done

Note: you may need to restart the kernel to use updated packages.

```
In [3]: import os

path = "."
```

```
print("Path to dataset files:", os.path.abspath(path))
```

Path to dataset files: c:\Users\huawei\Downloads

```
In [2]: import pandas as pd

df = pd.read_csv(r"C:\Users\huawei\Lab_1\IMDB Dataset.csv")

df.head()
```

```
Out[2]:
```

	review	sentiment
0	One of the other reviewers has mentioned that ...	positive
1	A wonderful little production. The...	positive
2	I thought this was a wonderful way to spend ti...	positive
3	Basically there's a family where a little boy ...	negative
4	Petter Mattei's "Love in the Time of Money" is...	positive

```
In [3]: df.describe()
```

```
Out[3]:
```

	review	sentiment
count	50000	50000
unique	49582	2
top	Loved today's show!!! It was a variety and not...	positive
freq	5	25000

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   review      50000 non-null  object
1   sentiment   50000 non-null  object
dtypes: object(2)
memory usage: 781.4+ KB
```

```
In [5]: # Check class distribution
print(df['sentiment'].value_counts())
```

```
sentiment
positive    25000
negative    25000
Name: count, dtype: int64
```

```
In [ ]:
```